HALCYON RIDGE HEALTH SYSTEM | ENTERPRISE AI ENABLEMENT ADVISORY

Enterprise AI Readiness & Current-State Assessment

HRHS is ready to pilot governed, low-risk AI in supply chain; it is not ready to scale AI broadly or clinically

Document	HRHS-AI-02 | Version 2.1 | November 2026 (supersedes v2.0) | Prepared by: Reginal Campbell, PMP
Audience and purpose	CIO/CTO, technology and transformation leadership, executive sponsors. Establish the evidence base for the strategic options in Document 03.
Decision status	Advisory work product. Findings are simulated case-study data [S]; estimates [E] and assumptions [A] are labeled. No decision has been made.
Classification	FICTIONAL CASE STUDY. All organizational details, interviews, figures, dates, and findings are simulated scenario data, estimates, or labeled assumptions. No real engagement, approval, or result is claimed. Not a description of any real organization.

Summary of conclusions

Across eight dimensions, HRHS rates 1 to 3 on a 5-point scale. Strengths are a modern cloud data platform for finance and supply chain, a functioning HIPAA privacy and security program, and strong clinical interest. The critical gaps are the absence of AI governance and lifecycle controls, constrained delivery capacity, and uneven data quality. These conditions favor an approach that builds minimum controls while delivering low-risk, data-ready use cases first, and they argue against both unmanaged vendor expansion and a large up-front platform build.

Assessment methodology

This is a simulated assessment. No interviews, document reviews, or data profiling were performed on a real organization. The scenario posits a six-week effort using four evidence sources: 31 stakeholder interviews across executive, IT, clinical, finance, supply chain, privacy, and security functions; review of 46 documents (policies, architecture diagrams, contracts, project portfolio, data dictionaries); data profiling of supply chain and revenue cycle extracts on the cloud data platform; and a walkthrough of both pilots. Each dimension was rated using the scale below. Ratings are analytical judgments [J] based on the simulated evidence [S].

Rating	Definition
1	Absent or ad hoc. No defined practice; outcomes depend on individuals.
2	Emerging. Some practice exists in pockets; not repeatable or enterprise-wide.
3	Defined. Documented and used in most relevant areas; gaps for AI-specific needs.
4	Managed. Consistently applied, measured, and adapted to AI requirements.
5	Optimized. Continuously improved; supports scaled AI with low friction.

Readiness matrix

Dimension	Rating	Key evidence	Finding
Strategy and use-case alignment	2	11 use cases from 7 sponsors via 4 intake routes; no shared prioritization criteria; pilots selected by departments without enterprise review.	F-01: Demand is real but unprioritized.
Data quality and accessibility	2 to 3 (varies by domain)	Supply: 112,000-item master, 18% of records incomplete or duplicate. Revenue cycle: 23% of denials coded "other." Clinical: data access mediated by EHR vendor.	F-02: Supply chain and revenue cycle data are usable with remediation; clinical data is not directly accessible.

Dimension	Rating	Key evidence	Finding
Architecture and integration	2	Three integration platforms (HL7 interface engine, cloud iPaaS, legacy ESB being retired). No approved pattern for AI model access, logging, or prompt and output retention.	F-03: No reference pattern for AI; risk of point-to-point sprawl.
Security, privacy, regulatory	3	Mature HIPAA program, SSO, BAAs standard. Vendor security questionnaire has no AI-specific items. P-01 data-use and retention terms not reviewed by Privacy.	F-04: Strong base controls; AI-specific gaps.
AI governance and model lifecycle	1	No AI inventory, risk tiering, evaluation standard, or monitoring practice. 2025 acceptable-use policy covers public generative AI tools only.	F-05: Governance absent; both pilots lack success criteria.
Technology operations and support	2	ITSM processes mature for applications; no runbooks, model performance monitoring, or AI incident categories.	F-06: AI cannot yet be supported in production.
Workforce readiness and change	2	High clinician interest; 14% vacancy in revenue cycle; frontline concern about job impact in billing; no AI literacy training.	F-07: Appetite is high; trust and capability are not yet built.
Delivery capacity and financial constraints	2	85,000 discretionary IT hours in FY2027, 92% committed; uncommitted hours are mostly generalist. Provisional AI envelope up to $4.5M (A-01).	F-08: Capacity, especially specialist skills, is the binding constraint.

Findings and supporting evidence

F-01. AI demand is broad but unprioritized

Sponsors described use cases in terms of tools rather than outcomes. Only 3 of 11 proposals stated a measurable business objective. Two sponsors were unaware of overlapping proposals in adjacent departments (for example, UC-05 and UC-06 both propose revenue cycle assistants from different vendors).

F-02. Data readiness varies sharply by domain

Profiling showed supply chain transaction and contract data are complete enough for spend classification, but item-master gaps would reduce accuracy for contract-utilization analysis until the top-spend categories are remediated. Two acquired hospitals remain on a legacy materials system until the planned FY2028 ERP migration. Revenue cycle claims and remittance data are available, but inconsistent denial reason coding limits root-cause analysis. Clinical documentation data cannot be evaluated independently of the EHR vendor.

F-03. Integration choices are being made one project at a time

Both pilots connect directly to vendor services outside the integration platforms. There is no common gateway to control which models are used, log activity, or apply data loss prevention. Each additional point solution increases the cost of future consolidation.

F-04. Baseline security and privacy are solid, but AI-specific controls are missing

HIPAA Privacy and Security Rule obligations are actively managed, and BAAs are standard practice. However, vendor assessments do not ask whether customer data is used to train models, how long prompts and outputs are retained, or how model changes are communicated. The P-01 agreement was signed by the clinical department using a standard template; Privacy has not confirmed retention and secondary-use terms (R-01).

F-05. AI governance and lifecycle controls do not yet exist

No body owns AI risk decisions. There is no inventory, no risk tiering, and no standard for evaluating outputs before or after deployment. P-01 reports clinician satisfaction anecdotally; no structured review of note accuracy has occurred. P-02 has run for five months without a defined success measure.

F-06. Operations cannot yet support AI in production

Service desk categories, escalation paths, and runbooks do not cover AI behavior such as degraded output quality, model version changes, or unexpected content. Incident response plans do not address AI-specific events.

F-07. Workforce appetite is high; capability and trust are not

Clinicians in P-01 report reduced after-hours documentation, but no measurement supports this yet. Revenue cycle staff expressed concern that AI is a prelude to reductions. No role-based training exists for reviewing or challenging AI outputs.

F-08. Capacity is the binding constraint, and skill mix matters more than total hours

[S] Of 85,000 discretionary IT hours for FY2027, about 78,200 are committed, leaving about 6,800, mostly application analysts, general data engineers, and project staff. [E] A governed first phase (Option B) requires about 14,000 IT hours: 9,400 generalist and 4,600 specialist (ERP item-master data 2,000; integration engineering 1,200; security architecture 900; enterprise architecture 500). The specialists are committed to ERP phase 2 and the EHR upgrade.

Capacity element [E]	Hours	Source
Total Option B delivery effort, year 1	14,000	Estimate from work breakdown
Generalist work covered by uncommitted internal capacity	6,800	Existing staff
Generalist work covered by additional contracted capacity	2,600	New contractors on AI work
Specialist work by internal specialists released from ERP/EHR	4,600	Requires release decision
Replacement contractors backfilling displaced ERP/EHR work	4,600	Not AI effort; avoids displacement
Total contracted hours (2,600 + 4,600) at ~$110/hour [A-04]	7,200 ($0.8M)	Counted once in one-time cost

[J] Funding the 7,200 contracted hours does not by itself make the plan executable. Specialists cannot be contracted in quickly because they know HRHS item-master data, integration patterns, and security architecture. They must be released by ERP and EHR leadership, and their replacements need 4 to 6 weeks to ramp up [A-06].

Initial application areas compared

Factor	Supply Chain Intelligence	Revenue Cycle Operations	Clinical Documentation Support
Business value	High. $640M spend; on-contract rate est. 71% vs 85% target (A-02).	High. $38M annual denial write-offs; initial denial rate 11.8%.	Moderate to high. Clinician burden and retention; financial value indirect.
Data readiness	Moderate. Usable with item-master remediation.	Moderate. Reason-code inconsistency.	Low for independent evaluation; vendor-mediated.
Integration complexity	Moderate. Cloud data platform and ERP.	Moderate to high. Revenue cycle vendor platform.	High. EHR-embedded workflow.
Compliance and safety exposure	Low. Minimal PHI; financial controls apply.	Moderate. PHI in claims; HIPAA applies; affects payment outcomes.	High. PHI; documentation enters the legal medical record; clinical safety review required.
Change management	Moderate. Small analyst and sourcing teams.	High. Large workforce, job-impact concerns.	High. Clinician workflow and accountability for notes.
Suitability as first enterprise use case	Strong candidate	Second wave after P-02 evaluation	Govern existing pilot; expand only after safety review

Critical gaps

ID	Gap	Linked findings and risk
GAP-01	No AI governance body, policy, inventory, or risk tiering	F-05; R-02 unmanaged AI adoption
GAP-02	No pre-deployment evaluation or production monitoring standard	F-05, F-06; R-04 undetected output errors
GAP-03	AI-specific vendor and privacy terms not assessed	F-04; R-01 PHI retention or secondary use
GAP-04	No reference integration pattern for AI services	F-03; R-06 architectural sprawl
GAP-05	Insufficient specialist delivery capacity for FY2027	F-08; R-05 capacity shortfall
GAP-06	No baselines for value measurement	F-01, F-02; R-07 unprovable benefits

Assumptions and known unknowns

ID	Assumption or unknown	How it will be resolved
A-01	CFO has indicated an FY2027 AI envelope of up to $4.5M, pending business case.	CFO confirmation before December 9
A-02	On-contract spend is 71% of $640M; derived from a 12-month extract with known item-master gaps.	Validate in Days 1 to 30 (WS-3)
A-03	Share of the $38M denial write-offs that is preventable is unknown; assumed 30% to 50% ($11.4M to $19.0M).	Root-cause sample in Days 31 to 60
A-04	Contracted capacity of 7,200 hours is obtainable at a blended ~$110 per hour (v1.0 used $85, judged too low).	Procurement market check
A-05	P-01 vendor will agree to amended data-use terms.	Privacy and Legal negotiation by January 31, 2027
A-06	Replacement contractors reach productivity in 4 to 6 weeks.	ERP and EHR program directors
A-07	AI-assisted workflow reduces preventable denial write-offs by 10% to 20%.	P-02 evaluation and Wave 2 pilot
A-08	The ERP phase 2 reporting module is not on the go-live critical path.	ERP program impact assessment
U-01	Whether the EHR vendor's native AI roadmap will overlap UC-07 and UC-08.	Vendor briefing request
U-02	Labor relations implications of revenue cycle automation.	HR review before any revenue cycle wave

Cross-functional dependencies

Supply chain use cases depend on the ERP team (item master, contract data) and Finance (savings validation). Revenue cycle use cases depend on the revenue cycle vendor and HR. Clinical use cases depend on the CMIO, nursing leadership, the EHR vendor, and Health Information Management. All use cases depend on Privacy, Security, and the Data Governance Council, which currently has stewards only in the finance domain.

Implications for the strategic options

Option	What the evidence suggests
A. Vendor-Led	Plausible for speed, but F-04 and F-05 mean each vendor deployment would carry its own unmanaged risk; F-03 sprawl would accelerate.
B. Governed, Phased	Matches the evidence: data-ready supply chain use cases can proceed while GAP-01 to GAP-04 are closed in parallel. Viable only if GAP-05 specialist capacity is committed.
C. Platform-First	Addresses GAP-03 and GAP-04 thoroughly, but F-08 makes it infeasible in FY2027 without deferring major committed programs, and GAP-06 means value would be unproven for 12 to 18 months.